




AI-race behavior has no universal safety ordering

Checkpoint and persona can move Unsafe play by most of the scale; group-size associations are smaller in the current N=3 -5 pilot

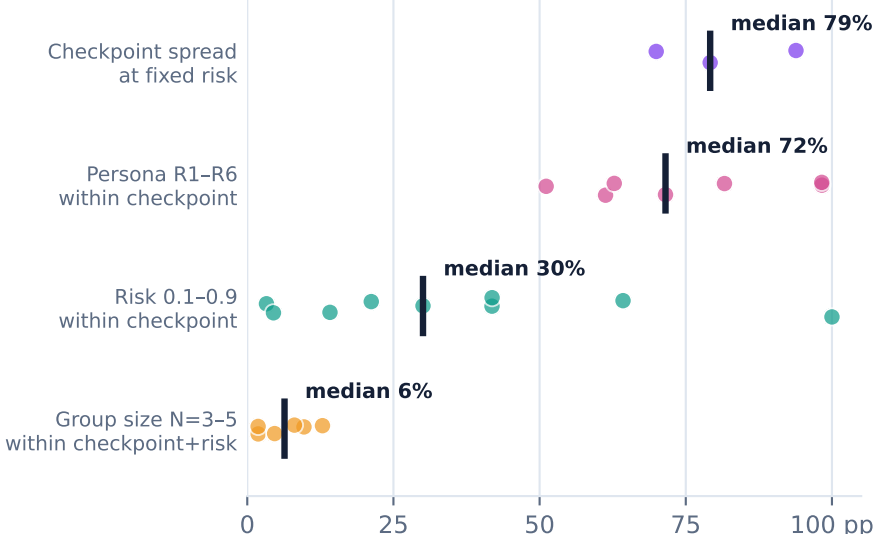
A Dyadic (N=2) risk fingerprints across benchmark layers**B Persona is a policy intervention, not cosmetic wording**



C N-player robustness



D Descriptive policy span



† EGT is a stationary population benchmark, not an LLM prediction. ‡ Human data are an external reference. All spans are descriptive; protocol-separated panels must not be read as one pooled causal model.