


Nine checkpoints, nine different risk responses -- and none matches the human pattern



Pilot, neutral/no-persona condition only; player-weighted rates across matched batches per checkpoint. The human bars are computed directly from the raw de-identified data (n=98/105/138 participants per risk cap), not taken as a scalar from the paper. Humans barely respond to the risk cap at all (64.2% -> 55.8% -> 55.8%); the checkpoints range from near-flat (GPT-5 nano, GPT-5.4 nano) to an almost perfect step function (Claude Opus 5: 100% -> 1% -> 0%).