


Human dynamics, refit directly on raw participant data, versus the one LLM checkpoint where the same fit converges



Logistic-regression coefficient on choosing Unsafe (95% CI)

Human: refit here on the paper's own de-identified data (N=2888 obs, 172 pair clusters), not a frozen scalar.
GPT-5.4 nano: the only one of 5 checkpoints where this specification converges (see INSIGHTS.md); others are not estimable.