


The position effect flips sign as risk-seeking framing increases (x markers = floor/ceiling or non-converged, not estimable)



N=3 races, OpenAI risk-persona sweep, cluster-robust by rep (10 clusters/cell). Negative = human-matching direction (behind $\rightarrow$ more Unsafe); positive = reversed (ahead $\rightarrow$ more Unsafe). Both models shown where estimable.