


Every checkpoint follows the same persona/framing gradient -- across three model families and three backend routes -- but each starts from a very different un-framed default

OpenAI-lineage checkpoints



Gemini and Claude



Assigned seat risk-attitude persona (none = un-framed baseline; R1 = most risk-averse ... R6 = most risk-seeking)