


The persona/grading gradient replicates in a newer model generation (GPT-5.6, run via a different backend entirely)



GPT-5.6 Luna/Terra: n=360 players/point, 0.1-0.6-0.9 risk pooled, results/frontier/bedrock_mantle/ (0 parse failures, 76/76 cells complete). Earlier three checkpoints: 2-player neutral-lane persona sweep (see persona_role_gradient.png). Pilot; never pooled across checkpoints.