

Pretrained SAE reconstruction fidelity

Normalized MSE at pre-action resid_post; n=600 decisions

Train Eval

Normalized MSE (lower is better)

0.0008

0.0006

0.0004

0.0002

0.0000

4

12

18

20

25

Transformer layer

